

cs109a_hw3

October 4, 2017

1 CS 109A/STAT 121A/AC 209A/CSCI E-109A: Homework 3

2 Multiple Linear Regression, Subset Selection, Cross Validation

Harvard University Fall 2017 Instructors: Pavlos Protopapas, Kevin Rader, Rahul Dave, Margo Levine

2.0.1 INSTRUCTIONS

- To submit your assignment follow the instructions given in canvas.
 - Restart the kernel and run the whole notebook again before you submit.
 - Do not include your name(s) in the notebook if you are submitting as a group.
 - If you submit individually and you have worked with someone, please include the name of your [one] partner below.
-

Your partner's name (if you submit separately): Everett Sussman, Jan Geffert

Enrollment Status (109A, 121A, 209A, or E109A): 109A

Import libraries:

```
In [2]: import sys
import numpy as np
import pandas as pd
import matplotlib
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import r2_score
from scipy import stats
import statsmodels.api as sm
from statsmodels.api import OLS
from sklearn.preprocessing import PolynomialFeatures
from sklearn.model_selection import KFold
from sklearn.linear_model import Ridge
from sklearn.linear_model import Lasso
from sklearn.linear_model import RidgeCV
from sklearn.linear_model import LassoCV
%matplotlib inline
sns.set_context("poster")
```

```
/Applications/anaconda/lib/python3.6/site-packages/statsmodels/compat/pandas.py:56:
FutureWarning: The pandas.core.datetools module is deprecated and will be removed in a
future version. Please use the pandas.tseries module instead.
from pandas.core import datetools
```

3 Forecasting Bike Sharing Usage

In this homework, we will focus on multiple linear regression and will explore techniques for subset selection. The specific task is to build a regression model for a bike share system that can predict the total number of bike rentals in a given day, based on attributes about the day. Such a demand forecasting model would be useful in planning the number of bikes that need to be available in the system on any given day, and also in monitoring traffic in the city. The data for this problem was collected from the Capital Bikeshare program in Washington D.C. over two years.

The data set is provided in the files `Bikeshare_train.csv` and `Bikeshare_test.csv`, as separate training and test sets. Each row in these files contains 10 attributes describing a day and its weather: - season (1 = spring, 2 = summer, 3 = fall, 4 = winter) - month (1 through 12, with 1 denoting Jan) - holiday (1 = the day is a holiday, 0 = otherwise) - day_of_week (0 through 6, with 0 denoting Sunday) - workingday (1 = the day is neither a holiday or weekend, 0 = otherwise) - weather - 1: Clear, Few clouds, Partly cloudy, Partly cloudy - 2: Mist + Cloudy, Mist + Broken clouds, Mist + Few clouds, Mist - 3: Light Snow, Light Rain + Thunderstorm + Scattered clouds, Light Rain + Scattered clouds - 4: Heavy Rain + Ice Pellets + Thunderstorm + Mist, Snow + Fog - temp (temperature in Celsius) - atemp (apparent temperature, or relative outdoor temperature, in Celsius) - humidity (relative humidity) - windspeed (wind speed)

and the last column 'count' contains the response variable, i.e. total number of bike rentals on the day.

3.1 Part (a): Data Exploration & Preprocessing

As a first step, identify important characteristics of the data using suitable visualizations when necessary. Some of the questions you may ask include (but are not limited to):

- How does the number of bike rentals vary between weekdays and weekends?
- How about bike rentals on holidays?
- What effect does the season have on the bike rentals on a given day?
- Is the number of bike rentals lower than average when there is rain or snow?
- How does temperature effect bike rentals?
- Do any of the numeric attributes have a clear non-linear dependence with number of the bike rentals?

```
In [3]: # Create Dataframes
trainDf = pd.read_csv("Bikeshare_train.csv")
testDf = pd.read_csv("Bikeshare_test.csv")
```

```
In [4]: trainDf.head()
```

```
Out [4]:
```

	Unnamed: 0	season	month	holiday	day_of_week	workingday	weather	temp	\
0	0	2	5	0	2	1	2	24	
1	1	4	12	0	2	1	1	15	
2	2	2	6	0	4	1	1	26	
3	3	4	12	0	0	0	1	0	
4	4	3	9	0	3	1	3	23	
		atemp	humidity	windspeed	count				
0	26	76.5833	0.118167	6073					

1	19	73.3750	0.174129	6606
2	28	56.9583	0.253733	7363
3	4	58.6250	0.169779	2431
4	23	91.7083	0.097021	1996

```
In [5]: # trainDf summary statistics
trainDf.shape
```

```
Out[5]: (331, 12)
```

```
In [6]: # testDf summary statistics
testDf.shape
```

```
Out[6]: (400, 12)
```

3.1.1 Importing the Data

We noticed that our data set appears to be very clean - all columns are filled with values, each data point is atomic, etc. We also noticed that our testing data table contains more data points than our training data table.

3.1.2 Number of Rentals per Day


The first question our group wanted to answer was what the overall distribution of our y variable, the number of bikes rented per day, looked like. Thus, we created the below histogram.

```
In [7]: # General histogram of count data
meanBikesRented = np.mean(trainDf["count"])
medianBikesRented = np.median(trainDf["count"])

fig, ax = plt.subplots(1)
ax = sns.distplot(trainDf["count"])
ax.set_title("Histogram of Number of Bikes Rented per Day")
ax.set_xlabel("Number of Bikes Rented")
ax.set_ylabel("Frequency")

ax.axvline(x = meanBikesRented, color="r", label = "Mean:
{0:.2f}".format(meanBikesRented))
ax.axvline(x = medianBikesRented, color="g", label = "Median:
{0:.2f}".format(medianBikesRented))
ax.legend()
```

```
Out[7]: <matplotlib.legend.Legend at 0x11f21e0f0>
```




The histogram above shows a fairly symmetrical distribution of values, as the mean and median are very similar. On an average day, 4598 bikes are rented. We also note that there tend to be small *sub-peaks* in our histogram at around 1500 bikes rented and 7500 bikes rented. There are no nonsensical values, such as negative bikes rented in a day.

3.1.3 Weekdays vs. Weekends

Next, we investigated whether bikes, on average, were more likely to be used during the weekend versus weekdays.

```
In [33]: dayGroup = trainDf.groupby("day_of_week", as_index=True)
ax = dayGroup["count"].mean().plot(kind="bar", title="Average number of Bikes Shared By
Day of Week")
plt.xticks(range(7), ["Sun", "Mon", "Tue", "Wed", "Thu", "Fri", "Sat"])
ax.set_ylabel("Mean Bikes Shared")
```

```
Out[33]: <matplotlib.text.Text at 0x122dc51d0>
```



```
In [9]: # Find mean for weekdays, mean for weekends
weekMean = np.mean(dayGroup["count"].mean()[1:-2])
weekendMean = np.mean([dayGroup["count"].mean()[0], dayGroup["count"].mean()[6]])
print("Mean bikes rented during the week: {}, Mean bikes rented during the weekend: {}".format(weekMean, weekendMean))
```

Mean bikes rented during the week: 4634.760765349032, Mean bikes rented during the weekend: 4503.6314465408805

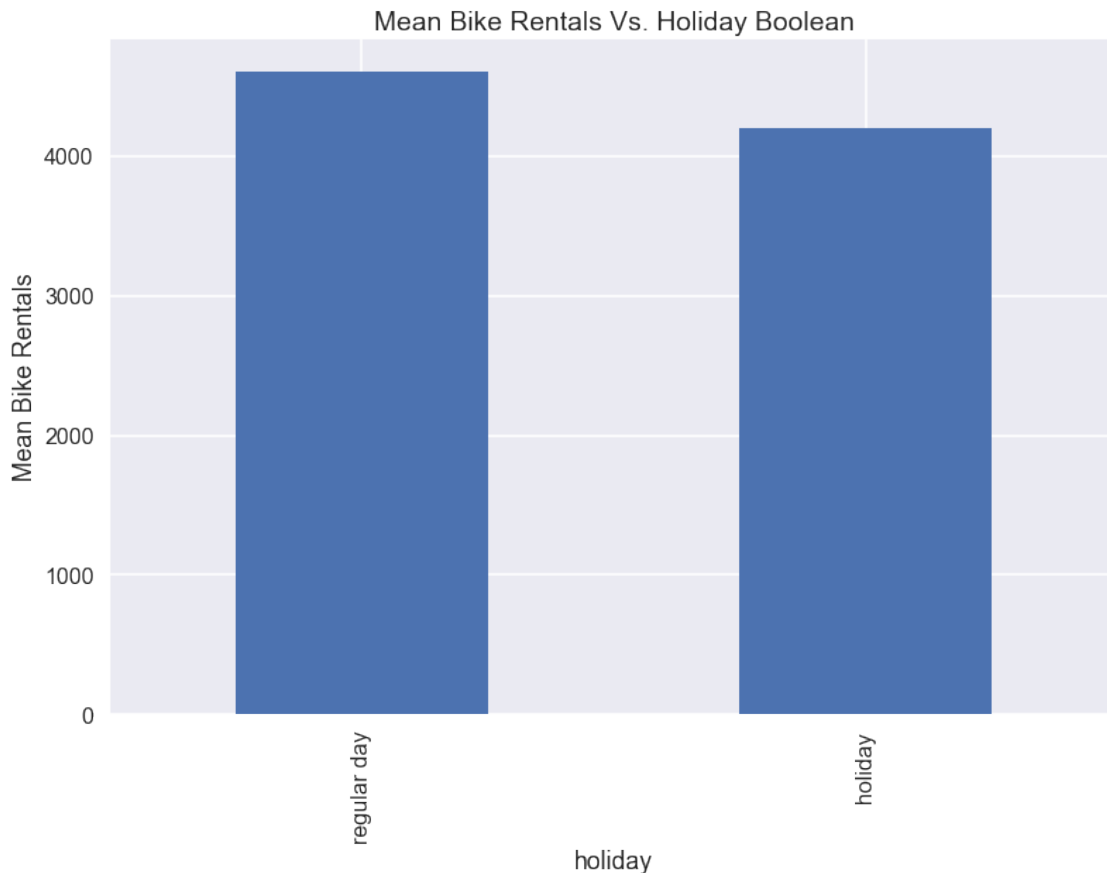
We find that, on average, slightly more bikes are rented during the week (4635 > 4504). This suggests the usage of rental bikes by commuters, as opposed to leisure.

3.1.4 Holidays

Next, we investigated whether the indicator of a day being a holiday impacted overall bike rentals.

```
In [10]: # Holiday Analysis
holidayGroup = trainDf.groupby("holiday")
ax = holidayGroup["count"].mean().plot(kind="bar", title="Mean Bike Rentals Vs. Holiday Boolean")
ax.set_ylabel("Mean Bike Rentals")
plt.xticks([0,1], ["regular day", "holiday"])
```

```
Out[10]: ([<matplotlib.axis.XTick at 0x11f517358>,
<matplotlib.axis.XTick at 0x11f517b00>],
<a list of 2 Text xticklabel objects>)
```



```
In [11]: hMeans = holidayGroup["count"].mean()
print("Mean Bike Rentals during a non-Holiday: {}, Mean Bike Rentals during a Holiday: {}"
      .format(hMeans[0], hMeans[1]))
```

```
Mean Bike Rentals during a non-Holiday: 4612.171875, Mean Bike Rentals during a
Holiday: 4199.181818181818.
```


We find that, on average, more bikes are rented during a non-holiday. This further hints at the commuter-hypothesis.

3.1.5 Seasonal effects

Now, our group wanted to determine whether the season impacted the relative popularity of bike rentals throughout the week. For example, in the summer, we would imagine that more people would rent a bike during the weekend than in the winter time, when the only people biking would be those that relied on bikes to commute.

To visualize this, in the below graph, we plotted 4 trend lines throughout the week detailing the normalized average bike rentals per day. To normalize each trend line, we divided the daily average by the seasonal average.

```
In [12]: seasonDayGroup = trainDf.groupby(["season", "day_of_week"])["count"].mean()
fig, ax = plt.subplots(1)
seasonNames = ["Spring", "Summer", "Fall", "Winter"]
for season in range(1,5):
    mean = np.mean(seasonDayGroup[season])
    normalizedRentals = list(map(lambda x: x / mean, seasonDayGroup[season]))
    ax.plot(normalizedRentals, label=seasonNames[season-1])
    ax.legend()
    ax.set_title("Normalized Average Bike Rentals Throughout the Week by Season")
    plt.xticks(range(7), ["Sun", "Mon", "Tue", "Wed", "Thu", "Fri", "Sat"])
    ax.set_ylabel("Normalized Average Bike Rental")
```



We find that there is considerably more variation in the spring time than in the other seasons. Also, in the spring, weekends are surprisingly unpopular in terms of average bike rentals compared to the weekdays. However, in contrast, in the summer, Fridays are very popular. Lastly, in the winter time, Saturday is the most popular day to rent a bike, relative to other days in the week. This was incredibly surprising, and throws a wrench into our initial hypothesis that most bike rentals were for commuting purposes only.

3.1.6 Weather Effects

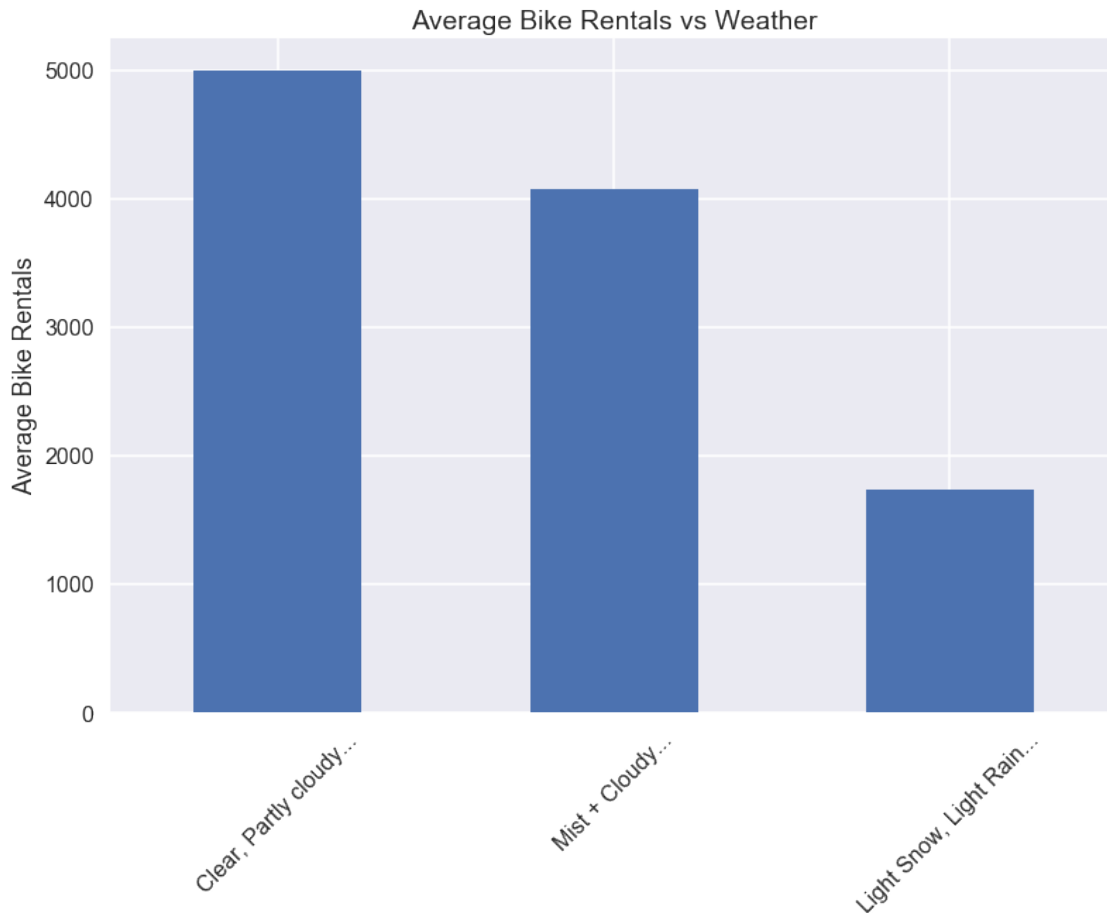
Now, our group investigated whether the severity of weather impacted bike rentals in any given day.

```
In [35]: weatherGroup = trainDf.groupby("weather")["count"].mean()

ax = weatherGroup.plot(kind="bar", title="Average Bike Rentals vs Weather")
plt.xticks(range(0,3), ["Clear, Partly cloudy...", "Mist + Cloudy...", "Light Snow,"
```

```
Light Rain..."], rotation=45)
ax.set_xlabel("")
ax.set_ylabel("Average Bike Rentals")
```

Out [35]: <matplotlib.text.Text at 0x1210110b8>



In the above graph, we see that, as expected, when the weather worsens, people rent fewer bikes.

```
In [14]: # Print out means:
weatherDiff = []

for weather in range(1,4):
    weatherMean = np.mean(weatherGroup[weather])
    weatherDiff.append(weatherMean - meanBikesRented)
    if weatherDiff[-1] > 0:
        print("Weather {} has a mean of: {}, which is {} more than the annual
mean.".format(weather, weatherMean, np.abs(weatherDiff[-1])))
    else:
        print("Weather {} has a mean of: {}, which is {} less than the annual
mean.".format(weather, weatherMean, np.abs(weatherDiff[-1])))
```

Weather 1 has a mean of: 5001.476415094339, which is 403.02928518497356 more than the annual mean.


Weather 2 has a mean of: 4077.1651376146788, which is 521.2819922946869 less than the

annual mean.

Weather 3 has a mean of: 1736.2, which is 2862.247129909366 less than the annual mean.

```
In [15]: fig, ax = plt.subplots(1)
         ax = sns.barplot(list(range(3)), weatherDiff)
         ax.set_title("Difference from Mean Bikes Rented Vs. Weather Category")
         plt.xticks(range(0,3), ["Clear, Partly cloudy...", "Mist + Cloudy...", "Light Snow,
         Light Rain..."], rotation=90)
         ax.set_ylabel("Difference in Bikes Rented from mean")
```

Out[15]: <matplotlib.text.Text at 0x11fec8198>



The above visualization takes this notion one step further, by revealing how far fewer bikes are rented in bad weather than the total average of bikes rented.

3.1.7 Temperature

Here, our group investigated the impact of temperature on bike rentals.


```
In [16]: fig, ax = plt.subplots(1)
         ax = sns.regplot(trainDf["temp"], trainDf["count"], fit_reg=False)
```

```

ax.set_title("Bikes Rented Vs. Temperature [°C]")
ax.set_xlabel("Temperature [°C]")
ax.set_ylabel("# of Bikes Rented")

```

Out[16]: <matplotlib.text.Text at 0x12019d550>



There seems to be a generally positive correlation of *number of bikes rented* and *temperature*, which makes sense, given our prior findings. At temperatures of 25 degrees or more, however, higher temperatures are associated with a reduction in ridership, presumably because riding becomes more cumbersome (and sweaty).



3.1.8 Checking for Non-linear Dependencies

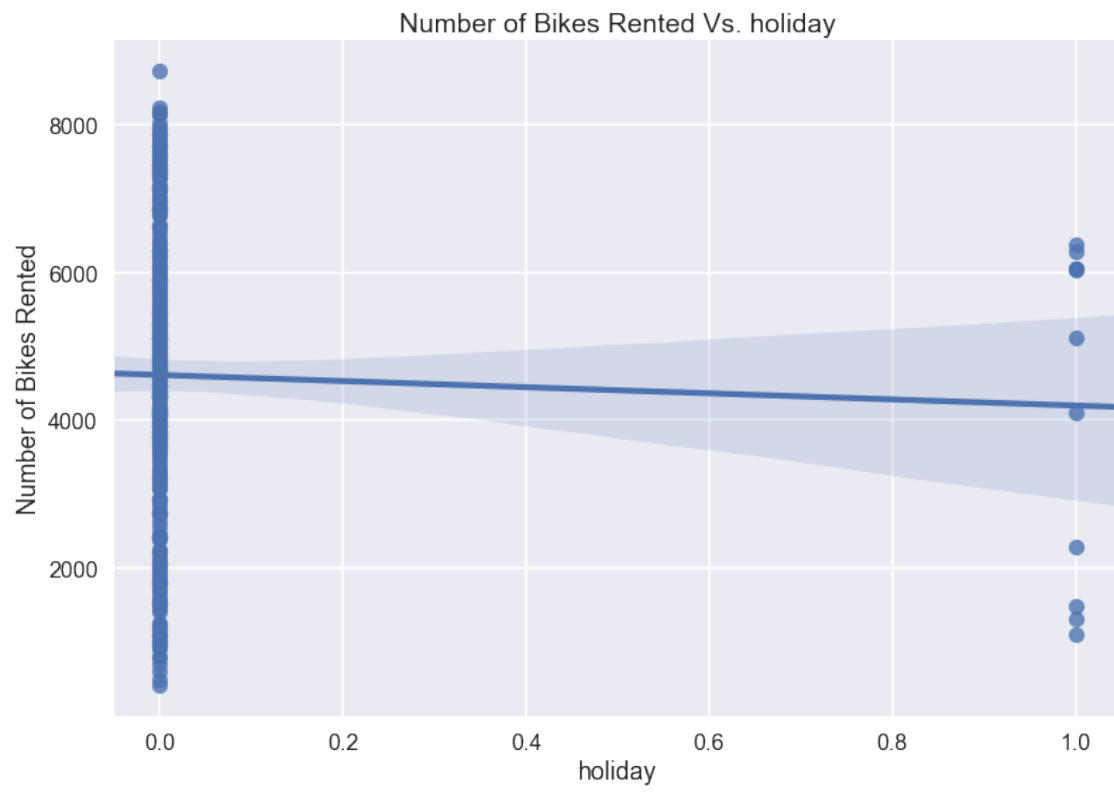
```


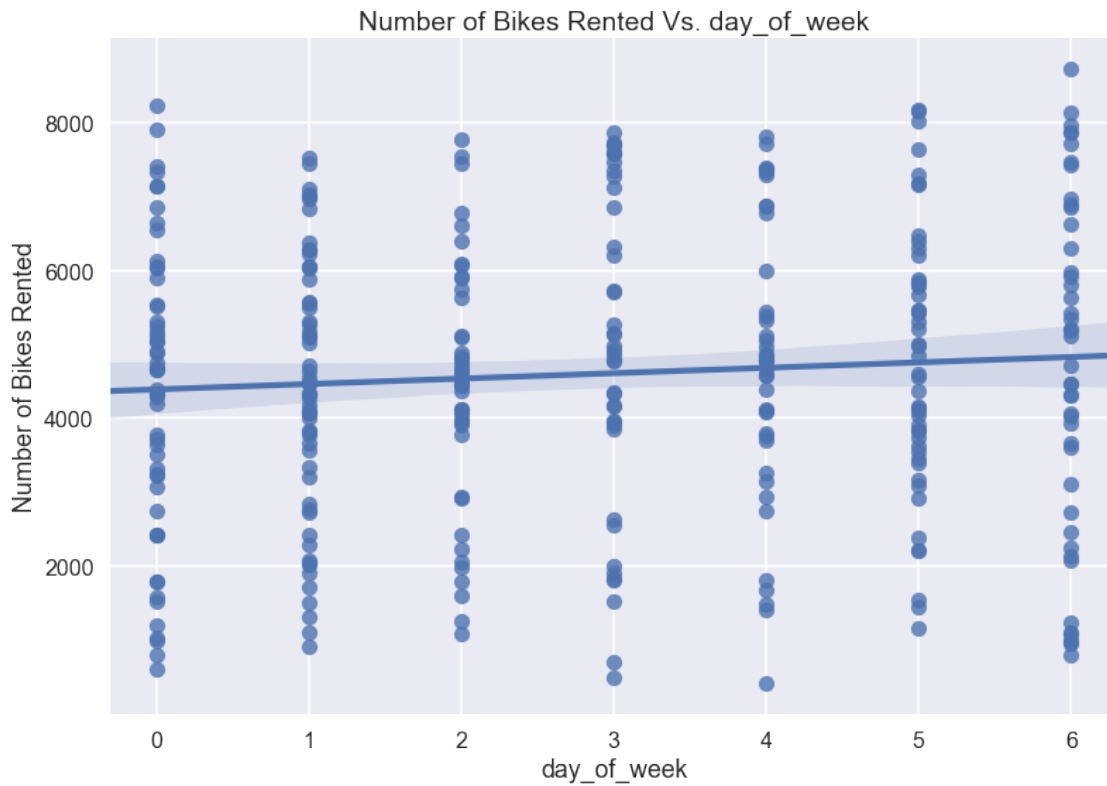
In [17]: initPredictors = ["season", "month", "holiday", "day_of_week", "workingday", "weather",
    "temp", "atemp", "humidity", "windspeed"]


for i, predictor in enumerate(initPredictors):
    fig, ax = plt.subplots()
    ax = sns.regplot(trainDf[predictor], trainDf["count"])
    ax.set_title("Number of Bikes Rented Vs. {}".format(predictor))
    ax.set_xlabel("{} {}".format(predictor))
    ax.set_ylabel("Number of Bikes Rented")

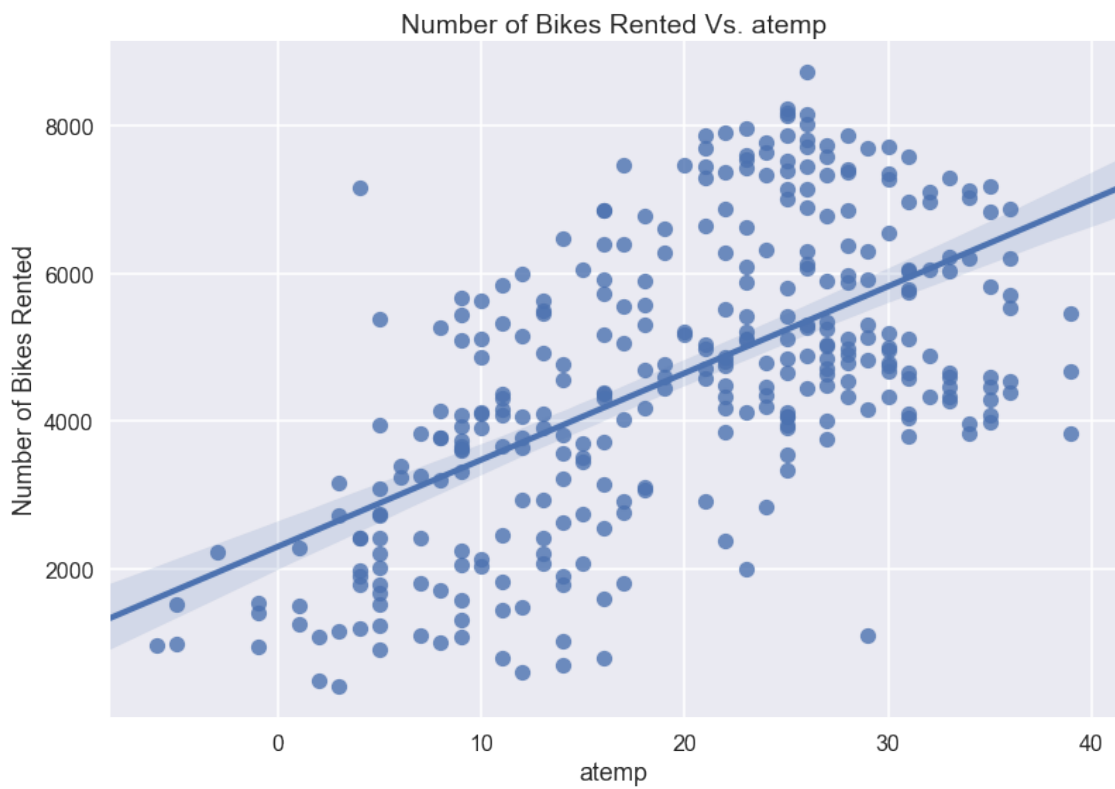

```

















We observe that season, month, temp, and atemp seem to have a non-linear relationship with number of bikes rented. The former is probably best described by a sinusoidal function due to the periodicity of seasons, whereas the latter two would fit a negative quadratic function.

We next require you to pre-process the categorical and numerical attributes in the data set:

- Notice that this data set contains categorical attributes with two or more categories. **Why can't they be directly used as predictors?** Convert these categorical attributes into multiple binary attributes using one-hot encoding: in the place of every categorical attribute x_j that has categories $1, \dots, K_j$, introduce $K_j - 1$ binary predictors $x_{j1}, \dots, x_{j,K_j-1}$ where x_{jk} is 1 whenever $x_j = k$ and 0 otherwise. ** Why is it okay to not have a binary column for the K_j -th category? **
- Since the attributes are in different scales, it is a good practice to standardize the continuous predictors, i.e. to scale each continuous predictor to have zero mean and a standard deviation of 1. This can be done by applying the following transform to each continuous-valued predictor j : $\hat{x}_{ij} = (x_{ij} - \bar{x}_j) / s_j$, where \bar{x}_j and s_j are the sample mean and sample standard deviation (SD) of predictor j in the training set. We emphasize that the mean and SD values used for standardization must be estimated using only the training set observations, while the transform is applied to both the training and test sets. ** Why shouldn't we include the test set observations in computing the mean and SD? **
- Provide a table of the summary statistics of the new attributes ('pd.describe()' function will help).

Hint: You may use the `pd.get_dummies` function to convert a categorical attribute in a data frame to one-hot encoding. This function creates K binary columns for an attribute with K categories. We suggest that you delete the last (or first) binary column generated by this function.

Note: We shall use the term "attribute" to refer to a categorical column in the data set, and the term "predictor" to refer to the individual binary columns resulting out of one-hot encoding.

3.1.9 Answers to Questions above

- We cannot use categorical variables directly as predictors since non-integer values have no meaning.
- The K_j -th category does not need to have a binary column as it is implied as the baseline, i.e. the 0-vector.
- As a general principle, we do not want to train on the test data. This way, we avoid fitting the training data which would lead to a less general model.

In [18]: # Create dummy variables

```
trainBinaryDf = pd.get_dummies(trainDf, columns=["season", "month", "day_of_week",
"weather"])
testBinaryDf = pd.get_dummies(testDf, columns=["season", "month", "day_of_week",
"weather"])

trainBinaryDf.head()
```

```
Out[18]:
```

	Unnamed: 0	holiday	workingday	temp	atemp	humidity	windspeed	count	\
0	0	0	1	24	26	76.5833	0.118167	6073	
1	1	0	1	15	19	73.3750	0.174129	6606	
2	2	0	1	26	28	56.9583	0.253733	7363	
3	3	0	0	0	4	58.6250	0.169779	2431	
4	4	0	1	23	23	91.7083	0.097021	1996	

	season_1	season_2	...	day_of_week_0	day_of_week_1	day_of_week_2	\
0	0	1	...	0	0	1	
1	0	0	...	0	0	1	
2	0	1	...	0	0	0	
3	0	0	...	1	0	0	
4	0	0	...	0	0	0	

	day_of_week_3	day_of_week_4	day_of_week_5	day_of_week_6	weather_1	\
0	0	0	0	0	0	
1	0	0	0	0	1	
2	0	1	0	0	1	
3	0	0	0	0	1	
4	1	0	0	0	0	

	weather_2	weather_3
0	1	0
1	0	0
2	0	0
3	0	0
4	0	1

[5 rows x 34 columns]

```
In [19]: # Normalize predictors
for column in ["temp", "atemp", "humidity", "windspeed"]:
    colMean = np.mean(trainBinaryDf[column])
    colSD = np.std(trainBinaryDf[column])
    trainBinaryDf[column + "_norm"] = [(x - colMean)/float(colSD) for x in
trainBinaryDf[column]]
    testBinaryDf[column + "_norm"] = [(x - colMean)/float(colSD) for x in
testBinaryDf[column]]
trainBinaryDf.head()
```

```
Out[19]:
```

	Unnamed: 0	holiday	workingday	temp	atemp	humidity	windspeed	count	\
0	0	0	1	24	26	76.5833	0.118167	6073	
1	1	0	1	15	19	73.3750	0.174129	6606	
2	2	0	1	26	28	56.9583	0.253733	7363	
3	3	0	0	0	4	58.6250	0.169779	2431	
4	4	0	1	23	23	91.7083	0.097021	1996	

	season_1	season_2	...	day_of_week_4	day_of_week_5	\
0	0	1	...	0	0	
1	0	0	...	0	0	
2	0	1	...	1	0	
3	0	0	...	0	0	
4	0	0	...	0	0	

	day_of_week_6	weather_1	weather_2	weather_3	temp_norm	atemp_norm	\
0	0	0	1	0	0.624743	0.651090	
1	0	1	0	0	-0.180583	-0.054841	
2	0	1	0	0	0.803704	0.852785	
3	0	1	0	0	-1.522794	-1.567551	
4	0	0	0	1	0.535262	0.348548	

	humidity_norm	windspeed_norm
0	0.922058	-0.930164
1	0.697907	-0.213825
2	-0.449062	0.805143
3	-0.332616	-0.269507
4	1.978781	-1.200843

[5 rows x 38 columns]

```
In [20]: trainBinaryDf.describe()
```

```
Out[20]:
```

	Unnamed: 0	holiday	workingday	temp	atemp	humidity	\
count	331.000000	331.000000	331.000000	331.000000	331.000000	331.000000	
mean	165.000000	0.033233	0.670695	17.018127	19.543807	63.385776	
std	95.695698	0.179515	0.470672	11.192515	9.930991	14.334789	
min	0.000000	0.000000	0.000000	-11.000000	-6.000000	25.416700	
25%	82.500000	0.000000	0.000000	7.500000	11.000000	52.702900	
50%	165.000000	0.000000	1.000000	18.000000	21.000000	63.291700	

75%	247.500000	0.000000	1.000000	26.000000	27.000000	73.500000
max	330.000000	1.000000	1.000000	38.000000	39.000000	97.250000

	windspeed	count	season_1	season_2	...	\
count	331.000000	331.000000	331.000000	331.000000	...	
mean	0.190833	4598.447130	0.217523	0.259819	...	
std	0.078240	1935.319338	0.413186	0.439199	...	
min	0.022392	431.000000	0.000000	0.000000	...	
25%	0.133083	3370.000000	0.000000	0.000000	...	
50%	0.178479	4648.000000	0.000000	0.000000	...	
75%	0.235380	5981.000000	0.000000	1.000000	...	
max	0.421642	8714.000000	1.000000	1.000000	...	

	day_of_week_4	day_of_week_5	day_of_week_6	weather_1	weather_2	\
count	331.000000	331.000000	331.000000	331.000000	331.000000	
mean	0.123867	0.145015	0.135952	0.640483	0.329305	
std	0.329929	0.352649	0.343256	0.480585	0.470672	
min	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	0.000000	0.000000	0.000000	
50%	0.000000	0.000000	0.000000	1.000000	0.000000	
75%	0.000000	0.000000	0.000000	1.000000	1.000000	
max	1.000000	1.000000	1.000000	1.000000	1.000000	

	weather_3	temp_norm	atemp_norm	humidity_norm	windspeed_norm
count	331.000000	3.310000e+02	3.310000e+02	3.310000e+02	3.310000e+02
mean	0.030211	-3.823729e-17	-1.214202e-16	-8.439037e-16	1.549616e-15
std	0.171428	1.001514e+00	1.001514e+00	1.001514e+00	1.001514e+00
min	0.000000	-2.507081e+00	-2.576025e+00	-2.652747e+00	-2.156128e+00
25%	0.000000	-8.516886e-01	-8.616201e-01	-7.463695e-01	-7.392325e-01
50%	0.000000	8.785869e-02	1.468532e-01	-6.572679e-03	-1.581428e-01
75%	0.000000	8.037042e-01	7.519372e-01	7.066402e-01	5.702098e-01
max	1.000000	1.877473e+00	1.962105e+00	2.365957e+00	2.954455e+00

[8 rows x 38 columns]

```
In [21]: # Drop k-th dummies
predictors = ["holiday", "workingday",
"temp_norm", "atemp_norm", "humidity_norm", "windspeed_norm",
"weather_1", "weather_2", "day_of_week_0", "day_of_week_1",
"day_of_week_2",
"day_of_week_3", "day_of_week_4", "day_of_week_5", "season_1",
"season_2", "season_3", "month_1", "month_2", "month_3", "month_4",
"month_5",
"month_6", "month_7", "month_8", "month_9", "month_10", "month_11",
"const"]

# Add a constant term to our predictors
trainBinaryDf = sm.add_constant(trainBinaryDf)
testBinaryDf = sm.add_constant(testBinaryDf)

trainPredict = trainBinaryDf[predictors]
testPredict = testBinaryDf[predictors]
```

3.2 Part (b): Multiple Linear Regression

We are now ready to fit a linear regression model and analyze its coefficients and residuals.

- Fit a multiple linear regression model to the training set, and report its R^2 score on the test set.
- *Statistical significance*: Using a t-test, find out which of estimated coefficients are statistically significant at a significance level of 5% (p-value<0.05). Based on the results of the test, answer the following questions:
 - Which among the predictors have a positive correlation with the number of bike rentals?
 - Does the day of a week have a relationship with bike rentals?
 - Does the month influence the bike rentals?
 - What effect does a holiday have on bike rentals?
 - Is there a difference in the coefficients assigned to temp and atemp? Give an explanation for your observation.
- *Residual plot*: Make a plot of residuals of the fitted model $e = y - \hat{y}$ as a function of the predicted value \hat{y} . Note that this is different from the residual plot for simple linear regression. Draw a horizontal line denoting the zero residual value on the Y-axis. Does the plot reveal a non-linear relationship between the predictors and response? What does the plot convey about the variance of the error terms?

```
In [22]: # Fit OLS multilinear regression on test set
X = trainPredict
Y = trainDf["count"]

allKModel = sm.OLS(Y, X).fit()
predictions = allKModel.predict(X)
allKModel.summary()
```

```
Out[22]: <class 'statsmodels.iolib.summary.Summary'>
"""
```

```

                        OLS Regression Results
=====
Dep. Variable:          count    R-squared:                0.576
Model:                  OLS      Adj. R-squared:           0.538
Method:                 Least Squares    F-statistic:        15.25
Date:                  Wed, 04 Oct 2017    Prob (F-statistic):   6.56e-42
Time:                  19:31:50    Log-Likelihood:       -2832.1
No. Observations:      331    AIC:                  5720.
Df Residuals:          303    BIC:                  5827.
Df Model:              27
Covariance Type:       nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
holiday	-616.6027	405.068	-1.522	0.129	-1413.706	180.500
workingday	-24.0933	174.143	-0.138	0.890	-366.777	318.590
temp_norm	924.3344	473.819	1.951	0.052	-8.058	1856.727
atemp_norm	311.9618	429.337	0.727	0.468	-532.898	1156.822

humidity_norm	-547.6638	113.029	-4.845	0.000	-770.085	-325.243
windspeed_norm	-254.7369	80.644	-3.159	0.002	-413.431	-96.043
weather_1	1581.9783	529.223	2.989	0.003	540.560	2623.396
weather_2	1565.4117	478.500	3.271	0.001	623.807	2507.016
day_of_week_0	-465.1450	269.154	-1.728	0.085	-994.794	64.504
day_of_week_1	-256.6501	172.675	-1.486	0.138	-596.445	83.145
day_of_week_2	-328.1845	204.810	-1.602	0.110	-731.214	74.845
day_of_week_3	37.6128	214.916	0.175	0.861	-385.304	460.530
day_of_week_4	-71.6425	208.337	-0.344	0.731	-481.612	338.327
day_of_week_5	-21.8317	200.742	-0.109	0.913	-416.858	373.194
season_1	-1226.1865	506.763	-2.420	0.016	-2223.407	-228.966
season_2	-327.3575	573.373	-0.571	0.568	-1455.654	800.939
season_3	-193.3050	449.171	-0.430	0.667	-1077.194	690.584
month_1	118.8358	505.353	0.235	0.814	-875.611	1113.282
month_2	207.7759	516.216	0.402	0.688	-808.047	1223.599
month_3	358.0167	511.391	0.700	0.484	-648.310	1364.344
month_4	452.1849	657.792	0.687	0.492	-842.234	1746.604
month_5	53.0233	700.991	0.076	0.940	-1326.403	1432.450
month_6	-673.4271	696.142	-0.967	0.334	-2043.313	696.458
month_7	-1161.1512	701.261	-1.656	0.099	-2541.109	218.806
month_8	-657.6397	684.628	-0.961	0.338	-2004.868	689.588
month_9	523.9804	548.284	0.956	0.340	-554.945	1602.906
month_10	605.0867	439.844	1.376	0.170	-260.449	1470.623
month_11	231.5175	413.966	0.559	0.576	-583.094	1046.129
const	3672.2940	664.433	5.527	0.000	2364.807	4979.781

Omnibus:	28.947	Durbin-Watson:	1.912
Prob(Omnibus):	0.000	Jarque-Bera (JB):	9.753
Skew:	0.054	Prob(JB):	0.00762
Kurtosis:	2.166	Cond. No.	1.05e+16

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
[2] The smallest eigenvalue is 7.45e-30. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.
"""

3.2.1 Assesment of The Regression Model Above

- The following predictors have a statistically significant correlation with the number of bike rentals:
 - *weather_1, weather_2, humidity_norm, windspeed_norm*
 - *month_4* April, *month_9* September, *month_10* October, *month_11* November, *season_1* Spring
 - *const*
- We determined these statistically significant predictors by looking at the $P > |t|$ values

above. If the value is less than .05, then that predictor's coefficient is significant.

- The day of a week does not have a statistically significant relationship with bike rentals.
- The month certainly influences the number of bike rentals (as evidenced by the significant predictors above)
- Ceteris paribus, a holiday does not have a significant relationship with bike rentals.
- There is no statistical significant difference in the coefficients assigned to *temp* and *atemp*, as both 95%-CIs have similar bounds, including 0. However, *temp* is almost significant, while *atemp* is nowhere close to being significant. This difference is surprising, but might be explained by the fact that apparent temperatures can vary significantly depending on other weather factors, while temperature is a baseline measure.

```
In [23]: testY = testDf["count"]
        testYHat = allKModel.predict(testPredict)
        print("Test R-Squared: {}".format(r2_score(testY, testYHat)))
```

Test R-Squared: 0.24934211146527574


3.3 Low R^2 Value Interpretation

The comparatively low R^2 value indicates that our model is overfitting the training data, especially given the larger R^2 value for the training data.

```
In [24]: residuals = list(map(lambda yhat, y: yhat - y, predictions, Y)) #YNorm instead of Y
        fig, ax = plt.subplots()

        ax.scatter(Y, residuals, alpha=.6)
        ax.set_title("Residual Plot from Multilinear Regression")
        ax.axhline(y=0, color="r")
        # ax.set_xlabel("Normalized True Count") True?
        ax.set_ylabel("Predicted Count")
```

Out[24]: <matplotlib.text.Text at 0x1229aa160>



3.3.1 Residuals analysis

The plot reveals a non-linear relationship between the predictors and response, since the error terms are not evenly distributed around the mean, and their variance depends on y , decreasing as y grows. This violates our regression assumption of normal, homoskedastic noise.

3.4 Part (c): Checking Collinearity

Does the data suffer from multi-collinearity? To answer this question, let us first analyze the correlation matrix for the data. Compute the (Pearson product-moment) correlation matrix for the predictor variables in the training set, and visualize the matrix using a heatmap. For categorical attributes, you should use each binary predictor resulting from one-hot encoding to compute their correlations. Are there predictors that fall into natural groups based on the correlation values?

Hint: You may use the `np.corrcoef` function to compute the correlation matrix for a data set (do not forget to transpose the data matrix). You may use `plt.pcolor` function to visualize the correlation matrix.

```
In [25]: dpal = sns.choose_colorbrewer_palette('diverging', as_cmap=True)
```




```
In [26]: plt.pcolor(X.ix[:, X.columns != "const"].corr(), cmap=dpal, alpha=.6)
         ax = plt.gca()
         ax.set_title("Correlation Matrix")
         labels = predictors[:-1]
         plt.xticks(range(len(labels)), labels, rotation=90)
         plt.yticks(range(len(labels)), labels)
```

/Applications/anaconda/lib/python3.6/site-packages/ipykernel_launcher.py:1:

DeprecationWarning:


.ix is deprecated. Please use
 .loc for label based indexing or
 .iloc for positional indexing

See the documentation here:

http://pandas.pydata.org/pandas-docs/stable/indexing.html#deprecate_ix

"""Entry point for launching an IPython kernel.

```
Out[26]: ([<matplotlib.axis.YTick at 0x122ce8b38>,
          <matplotlib.axis.YTick at 0x122cd4eb8>,
          <matplotlib.axis.YTick at 0x122df7e48>,
          <matplotlib.axis.YTick at 0x122de99e8>,
          <matplotlib.axis.YTick at 0x122dd9c88>,
          <matplotlib.axis.YTick at 0x122dc5cc0>,
          <matplotlib.axis.YTick at 0x122dc08d0>,
          <matplotlib.axis.YTick at 0x122dadcc0>,
          <matplotlib.axis.YTick at 0x122cefbe0>,
          <matplotlib.axis.YTick at 0x122e03550>,
          <matplotlib.axis.YTick at 0x122e18240>,
          <matplotlib.axis.YTick at 0x122e18cc0>,
          <matplotlib.axis.YTick at 0x122e1e780>,
          <matplotlib.axis.YTick at 0x122e27240>,
          <matplotlib.axis.YTick at 0x122e27cc0>,
          <matplotlib.axis.YTick at 0x122e2d780>,
          <matplotlib.axis.YTick at 0x122e33240>,
          <matplotlib.axis.YTick at 0x122e33cc0>,
          <matplotlib.axis.YTick at 0x122e38780>,
          <matplotlib.axis.YTick at 0x122e3f240>,
          <matplotlib.axis.YTick at 0x122e3fcc0>,
          <matplotlib.axis.YTick at 0x122e45780>,
          <matplotlib.axis.YTick at 0x122e4c240>,
          <matplotlib.axis.YTick at 0x122e4ccc0>,
          <matplotlib.axis.YTick at 0x122e51780>,
          <matplotlib.axis.YTick at 0x122e56240>,
          <matplotlib.axis.YTick at 0x122e56cc0>,
          <matplotlib.axis.YTick at 0x122e5f780>],
          <a list of 28 Text yticklabel objects>)
```



We observe the following relationships: * $temp_norm$ and $atemp_norm$ are unsurprisingly very much positively correlated. * $weather_1$ and $weather_2$ are negatively correlated, which makes sense since they derive from categorical data. * We can see "boxes" for $day_of_week_j$ and $month_k$ which is intuitive. * $temp_norm$ and $atemp_norm$ are correlated with the months. * Naturally, the seasons are correlated with the months.

Note that we excluded the constant term from our correlation matrix, as the variance of a constant is 0, so calculating the correlation of a variable with a constant would yield an undefined value.

3.5 Part (d): Subset Selection

Apply either one of the following subset selection methods discussed in class to choose a minimal subset of predictors that are related to the response variable: - Step-wise forward selection - Step-wise backward selection

We require you to implement both these methods *from scratch*. You may use the Bayesian Information Criterion (BIC) to choose the subset size in each method. Do these methods eliminate one or more of the redundant predictors (if any) identified in Part (c)? In each case, fit linear regression models using the identified subset of predictors to the training set. How do the test R^2 scores for the fitted models compare with the model fitted in Part (b) using all predictors?

```
In [27]: # Implement step-wise forward selection
         # The set of predictors
```

```

predictors = ["const", "holiday", "workingday", "temp_norm", "atemp_norm",
"humidity_norm", "windspeed_norm",
              "weather_1", "weather_2", "day_of_week_0", "day_of_week_1",
"day_of_week_2",
              "day_of_week_3", "day_of_week_4", "day_of_week_5", "season_1",
              "season_2", "season_3", "month_1", "month_2", "month_3", "month_4",
"month_5",
              "month_6", "month_7", "month_8", "month_9", "month_10", "month_11"]

forwardPredictors = []
forwardModels = []

for _ in range(len(predictors)):

    betterModel, betterPredictor, betterBIC = None, None, sys.maxsize

    for addedPredictor in predictors:
        X = trainBinaryDf[forwardPredictors + [addedPredictor]]
        newModel = sm.OLS(Y, X).fit()

        bestBIC = min([model.bic for model in forwardModels] + [betterBIC])
        if newModel.bic < bestBIC:
            betterModel = newModel
            betterPredictor = addedPredictor
            betterBIC = newModel.bic

    if betterModel:
        forwardModels.append(betterModel)
        forwardPredictors.append(betterPredictor)
    else:
        break

forwardBICs = [model.bic for model in forwardModels]
forwardRSquareds = [model.rsquared for model in forwardModels]
print("Predictors: {}".format(forwardPredictors))
print("RSquareds: {}".format(forwardRSquareds))
print("BICs: {}".format(forwardBICs))

```

```

Predictors: ['const', 'atemp_norm', 'humidity_norm', 'season_1', 'month_9',
'month_10', 'windspeed_norm', 'month_7']
RSquareds: [0.0, 0.36094811139379512, 0.41237745384733182, 0.46809945208979198,
0.48910179719539404, 0.5104124449264057, 0.5237587713442825, 0.53432886723575779]
BICs: [5954.1722009863015, 5811.7625734619323, 5789.7935136501037, 5762.6186708109581,
5755.0860658498423, 5746.7852518967693, 5743.438926291893, 5741.8118001623388]

```

3.6 Results

The step-wise forward regression yields a 6-parameter model. *temp_norm* and some *season/month* redundancies (by collinearity) have been resolved. We correctly see that the BIC of each iteration of this algorithm decreases, signifying an improvement overall in our model. Further, our R^2 values increase to just below the original model with all predictors. This is encouraging, since our new model is significantly smaller but about as powerful as the first model.

3.7 Testing

```

In [28]: testY = testDf["count"]
testPredict = testBinaryDf[forwardPredictors]
testYHat = forwardModels[-1].predict(testPredict)
print("Test R-Squared: {}".format(r2_score(testY, testYHat)))

```

Test R-Squared: 0.23390728606336764

Testing the resulting model on the test dataset yields a slightly lower R^2 value than the one in Part (b).

In [29]: # Implement step-wise backward selection

```
# The set of predictors
predictors = set(["const", "holiday", "workingday", "temp_norm", "atemp_norm",
"humidity_norm", "windspeed_norm",
"weather_1", "weather_2", "day_of_week_0", "day_of_week_1",
"day_of_week_2",
"day_of_week_3", "day_of_week_4", "day_of_week_5", "season_1",
"season_2", "season_3", "month_1", "month_2", "month_3", "month_4",
"month_5",
"month_6", "month_7", "month_8", "month_9", "month_10", "month_11"])
```

```
backwardPredictors = set(predictors)
backwardModels = []
```

```
for _ in range(len(predictors)):
```

```
    betterModel, betterPredictor, betterBIC = None, None, sys.maxsize
```

```
    for removedPredictor in backwardPredictors:
```

```
        X = trainBinaryDf[list(backwardPredictors - set([removedPredictor]))]
```

```
        newModel = sm.OLS(Y, X).fit()
```

```
        bestBIC = min([model.bic for model in backwardModels] + [betterBIC])
```

```
        if newModel.bic < bestBIC:
```

```
            betterModel = newModel
```

```
            betterPredictor = removedPredictor
```

```
            betterBIC = newModel.bic
```

```
    if betterModel:
```

```
        backwardModels.append(betterModel)
```

```
        backwardPredictors.remove(betterPredictor)
```

```
    else:
```

```
        break
```

```
backwardBICs = [model.bic for model in backwardModels]
```

```
backwardRSquareds = [model.rsquared for model in backwardModels]
```

```
print("Predictors: {}".format(backwardPredictors))
```

```
print("Rsquareds: {}".format(backwardRSquareds))
```

```
print("BICs: {}".format(backwardBICs))
```

```
Predictors: {'weather_2', 'day_of_week_5', 'month_8', 'humidity_norm', 'month_7',
'const', 'windspeed_norm', 'weather_1', 'workingday', 'month_6', 'day_of_week_2',
'season_1', 'day_of_week_3', 'day_of_week_4', 'holiday', 'temp_norm', 'day_of_week_1'}
Rsquareds: [0.57612017442132424, 0.57605080910137485, 0.57588543174157469,
0.5756513106142167, 0.57513980584466751, 0.57432856240861385, 0.57347456671855168,
0.57236259109017618, 0.57102272820283972, 0.56788327070409084, 0.56358808011277584,
0.5584752558959003]
BICs: [5820.9282263781215, 5815.1802696804134, 5809.507245152513, 5803.8877960781911,
5798.484420821208, 5793.3137231919627, 5788.175002389833, 5783.2346965933684,
5778.4680378881812, 5775.0795120710272, 5772.5512515638839, 5770.6044544609886]
```

3.7.1 Results of Backwards Selection

Here, we recognize several problems with the result of this backwards selection. First, the number of predictors in this model is still very large, and there are still many predictors that are colinear

with one another. Secondly, the R^2 values are decreasing, as well as our BIC values, which means that while our scoring metric (BIC) improves, we lose explanatory power of the data as well.

3.7.2 Testing

```
In [30]: testY = testDf["count"]
         testPredict = testBinaryDf[list(backwardPredictors)]
         testYHat = backwardModels[-1].predict(testPredict)
         print("Test R-Squared: {}".format(r2_score(testY, testYHat)))
```

Test R-Squared: 0.25606288780692477

Testing the resulting model on the test dataset yields a marginally higher R^2 value than the one in Part (b).

3.8 Part (e): Cross Validation

- Perform a 10-fold cross-validation procedure to select between the 3 competing models you have so far: the model with the best BIC from Step-wise forward selection, the model with the best BIC from Step-wise backward selection (if it is different), and the model with all possible predictors. Report the average R^2 across all 10 validation sets for each model and compare the results. Why do you think this is the case?
- Evaluate each of the 3 models on the provided left out test set by calculating R^2 . Do the results agree with the cross-validation? Why or why not?

```
In [31]: forwardModel = forwardModels[-1]
         backwardModel = backwardModels[-1]

         kf = KFold(n_splits=10)
         X = trainPredict
         Y = trainDf["count"]

         allKRSquares = []
         forwardRSquares = []
         backwardRSquares = []

         for trainIndices, validationIndices in kf.split(X):
             trainX = X.iloc[trainIndices]
             trainY = Y.iloc[trainIndices]

             validationX = X.iloc[validationIndices]
             validationY = Y.iloc[validationIndices]

             allKX = trainX
             allK = sm.OLS(trainY, allKX).fit()

             forwardX = trainX[list(forwardPredictors)]
             forward = sm.OLS(trainY, forwardX).fit()

             backwardX = trainX[list(backwardPredictors)]
             backward = sm.OLS(trainY, backwardX).fit()

             allKYHat = allK.predict(validationX)
             forwardYHat = forward.predict(validationX[list(forwardPredictors)])
             backwardYHat = backward.predict(validationX[list(backwardPredictors)])

             allKRSquares.append(r2_score(validationY, allKYHat))
             forwardRSquares.append(r2_score(validationY, forwardYHat))
             backwardRSquares.append(r2_score(validationY, backwardYHat))
```

```
print("Original model mean: {}, Forward model mean: {}, Backwards model mean: {}"
      .format(np.mean(allKRSquares), np.mean(forwardRSquares), np.mean(backwardRSquares)))
```

```
Original model mean: 0.34257855093958417, Forward model mean: 0.44036505151424993,
Backwards model mean: 0.4499215974237565.
```

3.8.1 Summary of 10-fold cross-validation

We see that our average R^2 score for 10-fold cross-validation from our original model was just .34, which is significantly smaller than both the forward and backwards models. This is likely due to the overfitting of our original model with the training dataset that it is given for each data split. As we have seen earlier, our step-wise backward model actually scored the best on our test data out of any other model. This is surprising for numerous reasons, but mostly because there still are a large amount of parameters in that model that might indicate overfitting.

In []: